

**Universität Stuttgart**

Institute of Parallel and  
Distributed Systems (IPVS)

Universitätsstraße 38  
D-70569 Stuttgart

# **Net-based Applications: Network Programming**

Adnan Tariq

(Based on the slides of Dr. Frank Dürr)

# Sockets

---

- Transport services
- Client/server model
- Sockets API
- Create, bind, close sockets
- Send and receive data
- Queues for connections
- Examples in C and Java



# Transport Services in the Internet

---

Transport Service offers inter-process communication

- End-to-end protocol connecting applications to each other
- Two classes of transport protocols: connection-less & connection-oriented

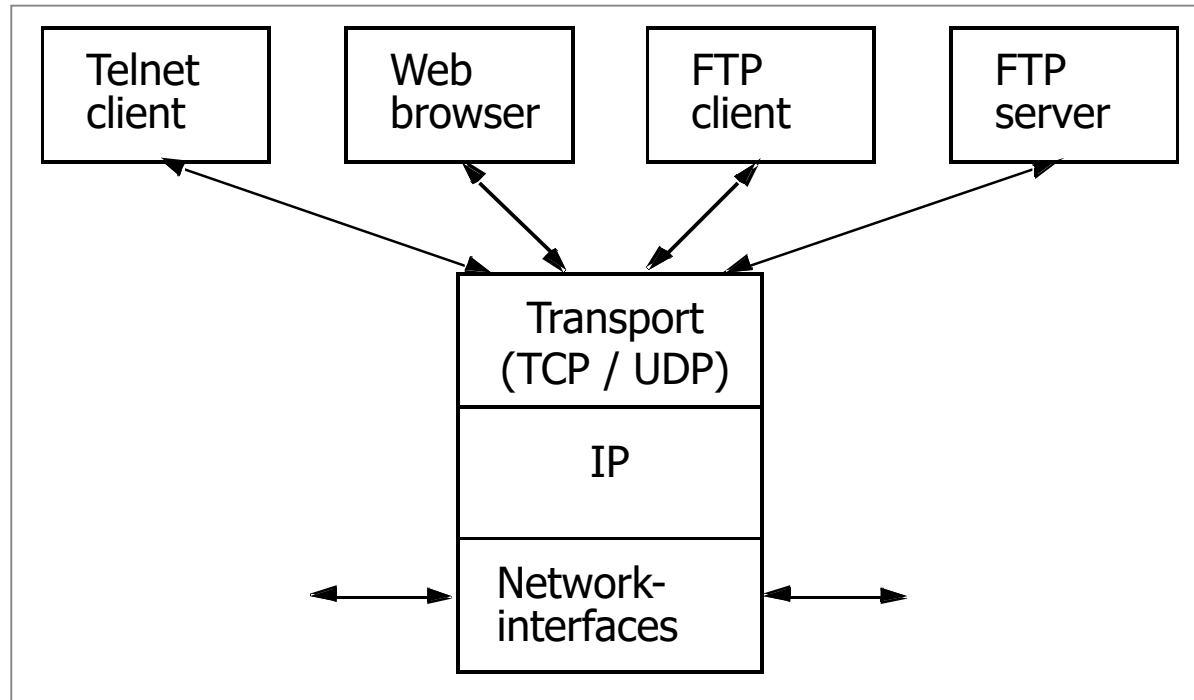
## Connection-less protocol: UDP

- No reliable communication, no order guarantees
  - Correctness of delivered messages guaranteed with checksums
- No flow control

## Connection-oriented protocol: TCP

- Reliable, ordered communication
- Byte-oriented communication
- Flow control
- Congestion control

# Addressing




- Network layer offers computer/computer communication and thus **addressing of computers (network interfaces) → IP addresses**
- Applications need inter-process communication and thus the **addressing of services/applications → ports**



# Client/Server Model

---



Server can serve several clients sequentially or in parallel  
➔ Server execution models



# Server Execution Models

---


There are three server execution models:

1. Independent requests with single-threaded server
2. Session-oriented processing with single-threaded server
3. Parallel processing with multi-threaded server




# 1. Independent Requests (Single-threaded Server)

- Requests are performed sequentially
- Typically no state is to be kept between requests
- Interleaved processing of requests from different clients
  - Server processes only one request at a time
  - To the clients, this looks like parallel processing of several clients




## 2. Session-oriented Processing (Single-threaded Server)

- Clients send requests within sessions
- Server executes requests sequentially
- Server executes all requests of single session in a row
- State for session can be maintained easily
- Reasonable and simple approach with connection-oriented transport protocol





# 3. Parallel Processing (Multi-threaded Server)



## Session-oriented Processing

- Primary server accepts session
- For each session a secondary server is created (or assigned)


## Independent requests

- All requests are sent to a single port
- Any idle server receives and processes request



# Sockets: A Transport Service API

- Problem:



- Interface between application- and protocol-software
- Sockets are an API (Application Program Interface)
- API is not standardized for TCP/IP:
  - Because: One interface may not suit all applications
  - Anyway: Sockets are widely used (de facto standard)



# API – Design Criteria

---

- In UNIX: **I/O-operations** usually follow the **open-read-write-close paradigm**
- This is not sufficient for complex interactions with network protocols, because:
  - ⇒ API must support both servers (passively waiting for requests) and clients (actively initiating transfers)
  - ⇒ To support datagrams: reasonable to give target address with datagram (not with open operation)
- API should not be restricted to TCP/IP



# Socket Abstraction (1)


---

Sockets are communication end-points

- Processes send and receive data via their local sockets

Connection-oriented paradigm:

- Two processes establish **bi-directional byte pipe** by connecting a pair of their sockets
- Sender writes a number of bytes to the corresponding local socket
- Receiver reads a number of bytes from the corresponding local socket



**IPVS**

Research Group  
Distributed Systems


Universität Stuttgart  
IPVS

# Socket Abstraction (2)

---

## Connection-less paradigm

- Sender sends datagrams through local sockets
  - Datagram includes destination address (IP-address, port number)
- Receiver receives datagrams through the local socket that is bound to the datagram's destination address



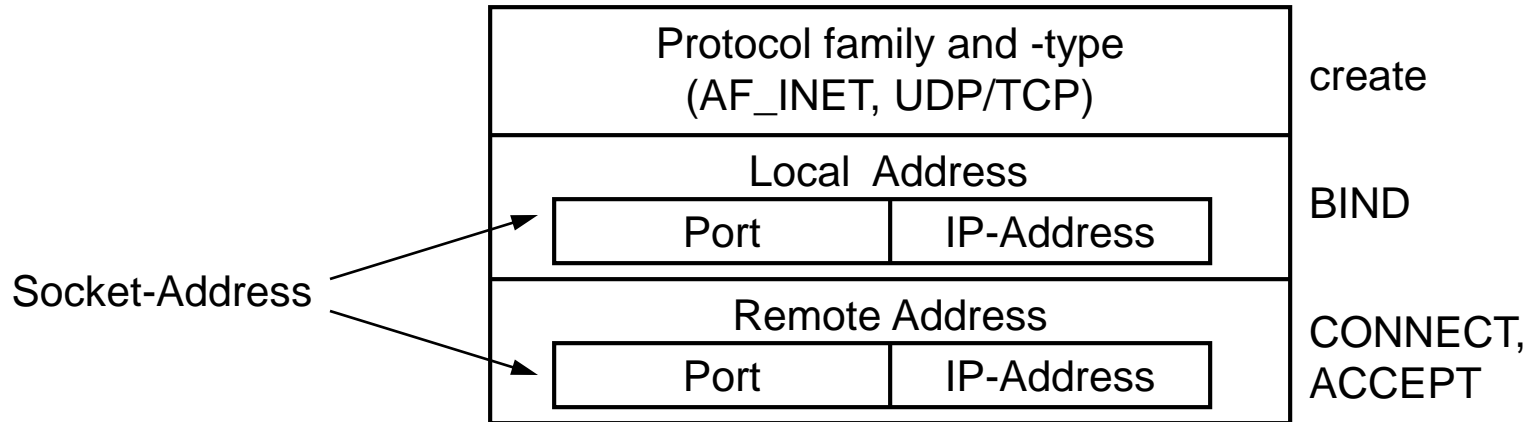
# Socket Primitives – Overview

- Socket access: similar to file access mechanism in UNIX:
  - In UNIX “everything is a file”
  - Access is identified via *file descriptors*. A *file descriptor* is an integer, which is connected to the “real” file → FIFO, Pipe, Terminal, ... , or network connection!
- Difference to files: sockets can be created without binding to a specific target address!

Files	Sockets	
	UDP	TCP
open	bind	connect, bind, accept
write	sendto	write
read	recvfrom	read
close	close	close, shutdown



# Socket – Data Structure



- Specify remote address with *connect* (can be omitted with UDP)
- Specify the recipient with every packet (*sendto*)
- Missing arguments are added by the OS
- Wildcards possible for IP-addresses
- Outgoing packets contain complete information
- Socket that is responsible for incoming packets is determined by information in packet and socket

# Socket Create and Close

---

`result = socket(pf, type, protocol)`

⇔ create signature for socket

`close(socket)`

⇔ close signature for socket

`pf:`

Protocol family (TCP/IP, AppleTalk, UNIX file system, ...)

`type:`

Type of communication

- Reliable and connection oriented service (`SOCK_STREAM`)
- Connection-less service (`SOCK_DGRAM`)
- *Raw* to permit privileged programs access to all protocol elements (`SOCK_RAW`)

`protocol:`

Specify protocol further (if *type* is not unique)

`result:`

Socket descriptor





# Bind Socket to Local Address

---

`bind(socket, localaddr, addrlen)`

`socket`:

socket descriptor

`localaddr`:

local address structure, e.g. in TCP/IP: port-number and IP-address

`addrlen`:

address length (Bytes)



# Connect Socket to Target Address

---

`connect(socket, destaddr, addrlen)`

`socket`:

socket descriptor

`destaddr`:

destination address structure, e.g. in TCP/IP: port-number and IP-address

`addrlen`:

address length (Bytes)

- Required for connection-oriented services
- Possible for connection-less services
  - No need to specify target address with every send operation

# Send Data

---

`write(socket, buffer, length)`

`socket`:

socket descriptor

`buffer`:

address pointing to data to be sent

`length`:

number of bytes to be sent

- *send*, *sendto*, *sendmsg*, *write*, and *writen* are possible sending operations
  - *send*, *write* and *writen* only for connected sockets

# Receive Data

---

`read(socket, buffer, length)`

**socket:**

socket descriptor

**buffer:**

address where to store received data

**length:**

maximal length of bytes to be received

- *read*, *readv*, *recv*, *recvfrom*, and *recvmsg* are possible receive operations

# Request Queue

---

## Problem in a client/server system:

What happens if the next client wants to connect before the current client is served?

## Solution:

To avoid refused connections, connection requests can be stored in a waiting queue.

`listen(socket, qlength)`

`socket:`

Socket Descriptor

`qlength:`

length of the waiting queue for this socket

# Accepting a Connection (1)

---

## Problem in a client/server system:

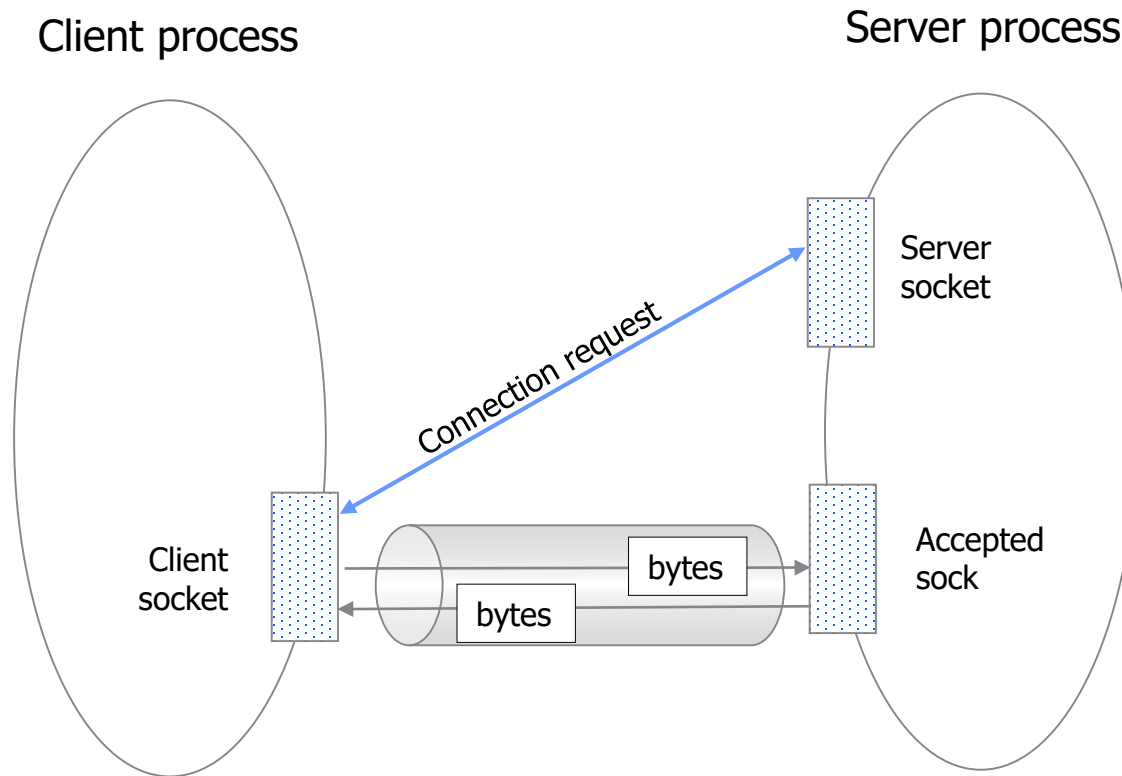
A server executes the operations *socket*, *bind*, and *listen*. How to establish a concrete connection?

## Solution:

*accept* operation: `newsock = accept(socket, addr, addrlen)`

- When a connection request comes from a client, *addr* and *addrlen* are set (by the system) according to the client address
- A new socket is created. This socket is connected to the client. It is now possible to communicate with the client through *newsock*
- The original socket stays open and continues to listen to connection requests.

# Accepting a Connection (2)



**IPVS**

Research Group  
Distributed Systems

Universität Stuttgart  
IPVS

# Sockets in Practice

---

There are further operations:

- Support of (single-threaded) servers that offer several services
  - Select readable socket from a set of server sockets
- Socket options (e.g. timeouts)
- ...

In practice:

- The operations described above do not necessarily exist exactly like that in every OS / are encapsulated in library functions
- Examples for C and Java following



# C Server (1)

---

```
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>

int main() {
    const int MAXBUF = 1024;
    const short SERVER_PORT = 4444;
    int serverSocketID, clientSocketID;
    struct sockaddr_in serverSocketAddr, clientSocketAddr;
    int addrLen = sizeof(struct sockaddr_in);
    char buffer[MAXBUF];
    int msgLen;

    serverSocketID = socket(PF_INET, SOCK_STREAM, 0);

    memset(&serverSocketAddr, 0, addrLen);
    serverSocketAddr.sin_family = AF_INET;
    serverSocketAddr.sin_addr.s_addr = INADDR_ANY;
    serverSocketAddr.sin_port = htons(SERVER_PORT);
    bind(serverSocketID, (struct sockaddr *)&serverSocketAddr, addrLen);
```



# C Server (2)

---

```
listen(serverSocketID, 5);
printf("Server running on port %d. Waiting for connections.\n",
      SERVER_PORT);
do {
    clientSocketID = accept(serverSocketID,
        (struct sockaddr *) &clientSocketAddr, &addrLen);
    printf("Accepting connection.\n");
    msgLen = recv(clientSocketID, buffer, MAXBUF, 0);
    buffer[msgLen] = '\0';
    printf("Received message: %s\n", buffer);

    /* this is the place where you could do something useful */

    send(clientSocketID, buffer, msgLen, 0);
    close(clientSocketID);
    printf("Closed connection. Waiting for new work.\n");
} while(1); // endless loop
}
```



# C Client (1)

---

```
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <netdb.h>

int main() {
    const int MAXBUF = 1024;
    const short PORT = 4444;
    char *message = "Hello World!";
    int clientSocketID;
    struct sockaddr_in serverSocketAddr;
    int addrLen = sizeof(struct sockaddr_in);
    struct hostent *serverInfo;
    char buffer[MAXBUF];
    int msgLen;

    clientSocketID = socket(PF_INET, SOCK_STREAM, 0);
```



## C Client (2)

---

```
memset(&serverSocketAddr,0,addrLen);
serverInfo = gethostbyname("taranis");
serverSocketAddr.sin_family = AF_INET;
serverSocketAddr.sin_addr = *((struct in_addr *) serverInfo->h_addr);
serverSocketAddr.sin_port = htons(PORT);

connect(clientSocketID, (struct sockaddr *) &serverSocketAddr,addrLen);

send(clientSocketID, message, strlen(message)+1, 0);

msgLen = recv(clientSocketID, buffer, MAXBUF, 0);
buffer[msgLen] = '\0';
printf("Received the following message from server: %s\n",buffer);

close(clientSocketID);
}
```



# C Server – Parallelism

---

```
do {
    clientSocketID = accept(serverSocketID, (struct sockaddr *)
        &clientSocketAddr, &addrLen);
    pid_t pid = fork();
    if (pid == -1) {
        // error
        perror("Could not fork");
    } else if (pid == 0) {
        // child process:
        msgLen = recv(clientSocketID, buffer, MAXBUF, 0);
        // this is the place where you could do something useful
        close(clientSocketID);
        exit(0); // child process exits after serving request
    } else {
        // fork returns child process id to parent process, i.e. != 0
        // parent process:
        close(clientSocketID); // close duplicate of client socket
    }
    // Avoid zombies calling wait, waitpid, ...
} while(1); // endless loop
```

# Protocol Independence

---

The **previous code is protocol dependent**: tailored to IPv4

- Using explicitly IPv4 definitions and 32 bit IPv4 addresses structures in code:
  - `struct sockaddr_in, PF_INET, AF_INET`
- New protocols required source code modifications to use new definitions and structures, e.g. IPv6 with 128 bit addresses
  - `struct sockaddr_in6, PF_INET6, AF_INET6`
- How can we write **protocol agnostic code**?
  - Compatible with IPv4, IPv6, etc. *without* code modifications
  - Application only specifies needs (“hints”) instead of concrete protocol
    - “I need reliable byte stream to host” (I don’t care which protocol)

→ **This is possible using function `getaddrinfo()`**



# User-friendly Name Strings – Lookup

- Resolve names using DNS, resolve services using file /etc/services
  - Replaces `gethostbyname()`

```
int getaddrinfo(const char *host, const char *service,  
               const struct addrinfo *hints, struct addrinfo **result);
```

- Parameters

- host: hostname or address string
- service: service name or decimal port number string
- hints: null pointer or pointer to `addrinfo` with type of information requested
- result: linked list of `addrinfo`, allocated by `getaddrinfo`, must be freed with `void freeaddrinfo(struct addrinfo *ai);`

```
struct addrinfo {  
    int ai_flags;        /* e.g., AI_PASSIVE */  
    int ai_family;       /* AF_XXX */  
    int ai_socktype;     /* SOCK_XXX */  
    int ai_protocol;     /* 0 or IPPROTO_XXX */  
    socklen_t ai_addrlen; /* length of  
                           ai_addr */  
    char *ai_canonname;  /* canonical name */  
    struct sockaddr *ai_addr; /* socket  
                               address */  
    struct addrinfo *ai_next; /* next list  
                               element */  
};
```

- Return value: 0 if OK, nonzero on error
- Results contain **all** address/port information for calling `socket`, `bind`, `connect`, `sendto`



# Protocol-independent Client Code

```
int sockfd; struct sockaddr_storage addr; socklen_t addrlen;

void init_socket(const char *hostname, const char *service) {
    struct addrinfo hints, *res, *ressave;
    bzero(&hints, sizeof(hints));
    hints.ai_family = AF_UNSPEC; /* any protocol will match our query */
    hints.ai_socktype = SOCK_STREAM /* reliable byte stream */
    if ( getaddrinfo(hostname, service, &hints, &res) != 0 ) handle_error();
    ressave = res;
    do { /* try to open a socket with each result entry until one works: */
        sockfd = socket(res->ai_family, res->ai_socktype, res->ai_protocol);
        if (sockfd < 0) continue; /* ignore this one */
        if (connect(sockfd, res->ai_addr, res->ai_addrlen) == 0) break;
        /* success */
        close(sockfd); /* ignore this one */
    } while ( (res = res->ai_next) != NULL );
    if (res == NULL) handle_error();
    freeaddrinfo(ressave);
    /* now you may use sockfd to communicate */
}
```

without any  
protocol  
specific  
code!



IPVS

Research Group  
Distributed Systems

Universität Stuttgart  
IPVS



# Protocol Independent Server Code

```
int sockfd; struct sockaddr_storage addr; socklen_t addrlen;

void init_socket(const char *hostname, const char *service) {
    struct addrinfo hints, *res, *ressave;
    bzero(&hints, sizeof(hints));
    hints.ai_flags = AI_PASSIVE; /* server passively waiting for conn. */
    hints.ai_family = AF_UNSPEC; /* any protocol will match our query */
    hints.ai_socktype = SOCK_STREAM /* reliable byte stream */
    if ( getaddrinfo(hostname, service, &hints, &res) != 0 ) handle_error();
    ressave = res;
    do { /* try to open a socket with each result entry until one works: */
        sockfd = socket(res->ai_family, res->ai_socktype, res->ai_protocol);
        if (sockfd < 0) continue; /* ignore this one */
        if (bind(sockfd, res->ai_addr, res->ai_addrlen) == 0) break;
        /* success */
        close(sockfd); /* ignore this one */
    } while ( (res = res->ai_next) != NULL );
    if (res == NULL) handle_error();
    freeaddrinfo(ressave);
    /* now you may use sockfd to communicate */
}
```

without any  
protocol  
specific  
code!



IPVS

Research Group  
Distributed Systems

Universität Stuttgart  
IPVS

# Java Client (1)

---

```
import java.io.*;
import java.net.*;
public class EchoClient {
    public static void main(String[] args) {
        Socket echoSocket = null;
        PrintWriter out = null;
        BufferedReader in = null;
        try {
            echoSocket = new Socket("taranis", 4444);
            out = new PrintWriter(echoSocket.getOutputStream(), true);
            in = new BufferedReader(new InputStreamReader(
                echoSocket.getInputStream()));
        } catch (UnknownHostException e) {
            System.err.println("Don't know about host: taranis.");
            System.exit(1);
        } catch (IOException e) {
            System.err.println("Couldn't get I/O for "
                + "the connection to: taranis.");
            System.exit(1);
        }
    }
}
```



# Java Client (2)

---

```
BufferedReader stdIn = new BufferedReader(  
    new InputStreamReader(System.in));  
String userInput;  
while ((userInput = stdIn.readLine()) != null) {  
    out.println(userInput);  
    System.out.println("echo: " + in.readLine());  
}  
out.close();  
in.close();  
stdIn.close();  
echoSocket.close();  
}  
}
```



# Java Client (Explanations)

---

```
echoSocket = new Socket("taranis", n)
```

⇒ Creates socket and connects to host *taranis* at Port *n*

```
out = new PrintWriter(echoSocket.getOutputStream(), true);
```

⇒ Get socket output stream and connect to new *PrintWriter* ("true" turns on "autoFlush")

```
in = new BufferedReader(new  
    InputStreamReader(echoSocket.getInputStream()));
```

⇒ Get socket input stream and open a new *BufferedReader*

To send data to server: write to *PrintWriter*

To receive data from server: read from *BufferedReader*



# Java Server (1)

---

```
import java.net.*;
import java.io.*;

public class EchoServer {
    public static void main(String[] args) {
        ServerSocket serverSocket = null;
        try {serverSocket = new ServerSocket(4444);}
        catch (IOException e) {
            System.err.println("Could not listen on port: 4444.");
            System.exit(1);
        }
        Socket clientSocket = null;
        try {clientSocket = serverSocket.accept();}
        catch (IOException e) {
            System.err.println("Accept failed.");
            System.exit(1);
        }
    }
}
```



# Java Server (2)

---

```
PrintWriter out =
    new PrintWriter(clientSocket.getOutputStream(), true);
BufferedReader in = new BufferedReader(new
    InputStreamReader(clientSocket.getInputStream()));
String inputLine, outputLine;
while ((inputLine = in.readLine()) != null) {
    outputLine = inputLine;
    out.println(outputLine);
    if (inputLine.equals("Bye."))
        break;
}
out.close(); in.close();
clientSocket.close(); serverSocket.close();
}
```

# Java Client & Server with UDP

---

## Client

```
// create socket and bind it
DatagramSocket clientSocket = new
    DatagramSocket(6001);
// send one datagram
byte[] data = new byte[1];
data[0] = 42;
DatagramPacket packet =
    new DatagramPacket(data,
        data.length);
// set receiver address
packet.setAddress(
    InetAddress.getByName(
        "localhost"));
packet.setPort(6000);
clientSocket.send(packet);
```

## Server

```
// create server socket and bind it
DatagramSocket serverSocket = new
    DatagramSocket(6000);
// allocate space for received
    datagram
byte[] data = new byte[256];
DatagramPacket packet = new
    DatagramPacket(data,
        data.length);
while (true) {
    // receive one datagram
    serverSocket.receive(packet);
    byte[] dataReceived =
        packet.getData();
    int bytesReceived =
        packet.getLength();
    // do something useful with data
}
```



# Java Server (Explanations)

---

```
serverSocket = new ServerSocket(4444)
```

⇒ Creates Server Object and listens to clients on port 4444

```
clientSocket = serverSocket.accept();
```

⇒ *accept*-method waits for client to connect to port 4444.

When a client connected successfully, *accept* returns a new socket

To send data to client: write to PrintWriter


To receive data from client: read from BufferedReader





# Further Literature

- Network programming in C:
  - W.R. Stevens, B. Fenner, A.M. Rudoff.  
UNIX Network Programming, Volume 1:  
The Sockets Networking API. 3rd Edition.  
Addison-Wesley. 2004.
- Network programming in Java:
  - Sun Microsystems: *Java Tutorial – Custom Networking Trail*.  
Available online:  
<http://docs.oracle.com/javase/tutorial/networking/>



IPVS

Research Group  
Distributed Systems

Universität Stuttgart  
IPVS

# Summary

---

- Sockets are an API for the transport layer
  - Transfer byte stream or datagrams (connection-oriented and –less transport services)
- Based on the “file philosophy” in Unix
  - Socket is a file descriptor
- Specific functions for creating, binding, accepting, connecting sockets
- Programming is “low-level”
  - Application has to parse and interpret transmitted data (bytes)
- Often used as basis for “higher” interactions like RPC

